

Project 3: Demo

Christopher Di Giacomo

40133600

10 November 2023

COMP479 2232 D

Dr. Sabine Bergler

Table of Contents

ENVIRONMENT/SETUP	3
SUBPROJECT1: FIRST-RUN WALKTHROUGH	3
MAIN.....	3
GETTOKENS().....	4
TOKENIZER()	5
SPIMI()	5
SAVE2JSON()	6
NAÏVE()	7
GETSTATS().....	8
SUBPROJECT 2: FIRST-RUN WALKTHROUGH.....	9
MAIN.....	9
QUERYMANAGER()	9
GETOPERATION().....	9
QUERYTEST()	11
CONJUNCTION()	12
INTERSECT()	12
DISJUNCTION().....	13
QUERYTERMRANK().....	13
CONVERTPOSTINGSLISTS().....	14
BM25()	14

Environment/Setup

Directories:

- **/reports:** This directory contains the report and demo for this project
- **/src:** This directory contains the *subProject1.py* and *subProject2.py* files
- **/indexes:** This directory contains the indexes in JSON format. It is (re)generated at runtime for either *.py* files.
- **/Corpus:** This directory contains all of the reuters collection *.sgm* files

Notes:

- The current working directory must be the parent directory of the submission folder. Although the */reports* subdirectory will be present in the project folder structure, within the text editor of choice, this will not interfere with any computation.
- If using an extension based IDE like VS Code, you must run the *pipeline.py* file as a Python file, and not via any code runner extensions.

Subproject1: First-Run Walkthrough

Running *subProject1.py* will trigger the *main()* function call and result in the following sequence of operations until terminal output.

Main

The main function below executes, retrieving the collection, building both indexes, and calling *getStats()* to display the performance data.

```
def main():
    """
    Main function to execute the token retrieval, index building, and statistics generation process.
    It retrieves tokens, builds indexes using both naive and SPIMI approaches, and then prints statistics.
    """

    # Measure time taken to create token streams
    start = time.time()
    tokenStream = getTokens()
    end = time.time()
    tokensTime = round(end - start, 3)

    # Build SPIMI and Naive indexes
    spimiIndex, spimiTime = spimi(tokenStream, 10000)
    naiveIndex, naiveTime = naive(tokenStream, 10000)

    # Generate and print statistics
    getStats(naiveIndex, naiveTime, spimiIndex, spimiTime, tokensTime)
```

getTokens()

The first function called in the main is the *getTokens()* function, displayed below. As in the previous report, it iterates through all .sgm files of the collection and calls the *tokenizer()* function on each document. The resulting tokenized collection is returned, in variable *tokensList*.

```
26 def getTokens():
27     """
28     Retrieves and tokenizes documents from a specified corpus. Each .sgm file in the
29     corpus is parsed, and documents within are tokenized. Tokens are associated with their
30     document ID.
31
32     Returns:
33     | list: A list where each element represents the tokens of a document, indexed by document ID.
34     """
35
36     tokensList = []
37     # Inform about the start of token retrieval process
38     print(f'Retrieving tokens ...')
39     # Loop through each file in the corpus
40     for file_number in range(22):
41         with open(f'./Corpus/reut2-{str(file_number).zfill(3)}.sgm', 'r', encoding='windows-1252') as f:
42             # Parse the file content using BeautifulSoup
43             soup = BeautifulSoup(f, 'html.parser')
44             documents = soup.find_all('text')
45
46             for document in documents:
47                 doc_type = document.get('type')
48                 text_parts = []
49
50                 # Different handling for 'BRIEF' type documents
51                 if doc_type == 'BRIEF':
52                     title = document.title.get_text() if document.title else ""
53                     text_parts.append(title)
54                 else:
55                     title = document.title.get_text() if document.title else ""
56                     body = document.body.get_text() if document.body else ""
57                     text_parts.append(title)
58                     text_parts.append(body)
59
60                 # Combine title and body for tokenization
61                 full_text = ' '.join(text_parts)
62                 tokens = tokenizer(full_text)
63                 tokensList.append(tokens)
64
65     # Indicate completion of token retrieval
66     print(f'Tokens retrieved!')
67     return tokensList
```

tokenizer()

The tokenizer() function is displayed below. As in project 2, it returns the tokenized inputted string with respect to a regex pattern, using NLTK's `regex_tokenize()`.

```
4 def tokenizer(text):
5     """
6     This function takes in a text and returns a list of tokens. It uses a regular expression to
7     handle acronyms, abbreviations, and general word patterns including words with hyphens or apostrophes.
8
9     Args:
10         text (str): The text to be tokenized.
11
12     Returns:
13         list: A list of tokens extracted from the input text.
14     """
15
16     # Pattern captures abbreviations/acronyms (e.g., U.S.A.) and general word forms.
17     pattern = r'\b(?:[A-Z]{1,2}\.)*[A-Z]{1,2}\.?\b|\b\w+(?:[-\']\w+)*\b'
18     return nltk.regex_tokenize(text, pattern)
```

As mentioned, once the `getTokens()` function has tokenized all .sgm files using `tokenizer()`, it returns a list of tokenized documents which are in the form of lists of tokens. This is returned to the main method, which has calculated the elapsed time for this phase of subproject 1. Next, it begins by constructing the `spimiIndex`, whose implementation can be found below.

spimi()

The function builds the inverted index in a single pass using the `tokensList`, returned by `getTokens()`. It also has an optional parameter for subproject 1, which halts the construction

```
69 def spimi(tokensList, testCount=None):
70     """
71     Builds an inverted index using the Single-Pass In-Memory Indexing (SPIMI) algorithm.
72     It processes tokenized documents, generating a dictionary where each term is associated
73     with a list of tuples containing document IDs and term frequency.
74
75     Args:
76         tokensList (list): A list of lists, each containing tokens from a document.
77         testCount (int, optional): A threshold for the number of term-document pairs to process.
78
79     Returns:
80         tuple: A tuple containing the inverted index and the time taken to build it.
81     """
82     index = {}
83     docID = 1
84     pairCount = 0
85     # Print collection size and start building index
86     print(f'Collection size: {len(tokensList)}')
87     print(f'Building SPIMI index ...')
88     startTime = time.time()
89     for tokens in tokensList:
90         for token in tokens:
91             pairCount += 1
92             # Update or add token in the index
93             if token in index:
94                 postingsList = index[token]
95                 if postingsList[-1][0] == docID:
96                     postingsList[-1] = (docID, postingsList[-1][1] + 1)
97             else:
98                 postingsList.append((docID, 1))
99             index[token] = [(docID, 1)]
100
101             if testCount and pairCount >= testCount:
102                 break
103         docID += 1
104         if testCount and pairCount >= testCount:
105             break
106
107     # Sort index by keys (terms) before saving
108     index = dict(sorted(index.items(), key=lambda item: item[0]))
109     endTime = time.time()
110
111     print(f'SPIMI index built!')
112     save2json(index, 'spimi.json')
113     return index, endTime - startTime
```

after the first *testCount* terms have been processed. In *main()* we set *testCount* to 10 000. The function builds the index in one pass over the data within *tokensList*. It adds terms to the index, along with a postings list of the *docIDs* it occurs in. The postings are in the form of *docID-term frequency* tuples, and everytime a term reoccurs in the same document, the term frequency for that document is incremented. Throughout the construction process the function is being timed, and when it returns the final index is also returns the elapsed time.

`save2json()`

Just before returning, the function calls *save2json()* which is seen below. It simply outputs the dictionary structure to a JSON file object within the */indexes* directory.

```
115 def save2json(data, filename):
116     """
117     Saves a given data object to a JSON file.
118
119     Args:
120     | data: The data to be saved.
121     | filename (str): The name of the file to save the data in.
122     """
123
124     # Create directory if it doesn't exist
125     if not os.path.exists('indexes'):
126         os.makedirs('indexes')
127
128     # Save data to a JSON file
129     with open(f'./indexes/{filename}', 'w') as f:
130         json.dump(data, f, indent=4)
```

naïve()

Back in main(), the spimi index and its construction time is retrieved is retrieved, and naïve() is now called, which can be seen below. It has the same parameters as spimi(), which work the same way. However, this function stores term-docID pairs in a list F, which it then sorts and

```
132 def naïve(tokensList, testCount = None):
133     """
134     Builds an inverted index using a naive approach. It processes term-document pairs,
135     sorts them, and then groups them to create postings lists.
136
137     Args:
138         tokensList (list): A list of tokenized documents.
139         testCount (int, optional): A threshold for the number of term-document pairs to process.
140
141     Returns:
142         tuple: A tuple containing the inverted index and the time taken to build it.
143     """
144
145     F = [] # List to store term-docID pairs
146     docID = 1
147     print(f'Building Naive index ...')
148
149     startTime = time.time()
150     # Create term-docID pairs
151     for tokens in tokensList:
152         for token in tokens:
153             F.append((token, docID))
154             if testCount and len(F) == testCount:
155                 break
156         docID += 1
157         if testCount and len(F) == testCount:
158             break
159
160     # Sort and remove exact duplicates
161     F.sort()
162     F = list(dict.fromkeys(F))
163
164     # Create postings lists
165     index = {}
166     for term, docID in F:
167         if term in index:
168             if index[term][-1] != docID:
169                 index[term].append(docID)
170         else:
171             index[term] = [docID]
172
173     endTime = time.time()
174     print(f'Naive index built!')
175     save2json(index, 'naive.json')
176     return index, endTime - startTime
```

cleans for duplicates, all before actually constructing the index. When, building the index, it works exactly as does spimi(), although it does not maintain term frequency. It also writes the index to a JSON file, and returns both the index and the elapsed time.

Back in `main()`, both indexes are now built, and both of their elapsed times have been retrieved.

`getStats()`

The next function call is to `getStats()` which can be seen below. This function simply takes all of the previously gathered data as input parameters, and displays it in a proper manner via

```
178 def getStats(naiveIndex, naiveTime, spimiIndex, spimiTime, tokensTime):
179     """
180     Prints statistics comparing the Naive and SPIMI indexing methods. It displays the time taken
181     to create token streams, build indexes, and the sizes of the created indexes.
182
183     Args:
184         naiveIndex (dict): The inverted index created using the naive approach.
185         naiveTime (float): The time taken to build the naive index.
186         spimiIndex (dict): The inverted index created using the SPIMI approach.
187         spimiTime (float): The time taken to build the SPIMI index.
188         tokensTime (float): The time taken to create token streams.
189     """
190     print('\n===== STATISTICS =====\n')
191     # Token streams creation time
192     print(f'Token streams creation time (sec): {tokensTime}')
193     # SPIMI Index Construction Statistics
194     print('\n-----\n| SPIMI Construction |\n-----')
195     print(f'\nTime taken for SPIMI index to process 10 000 terms (ms): {round(spimiTime * 1000, 3)}')
196     print(f'\nSize of SPIMI index: {len(spimiIndex)} terms')
197     # Naive Index Construction Statistics
198     print('\n-----\n| Naive Construction |\n-----')
199     print(f'\nTime taken for Naive index to process 10 000 terms (ms): {round(naiveTime * 1000, 3)}')
200     print(f'\nSize of Naive index: {len(naiveIndex)} terms')
201     # Differences between SPIMI and Naive Methods
202     print('\n-----\n| Differences |\n-----')
203     time_diff = naiveTime - spimiTime
204     time_diff_percent = (time_diff / naiveTime) * 100
205     print(f'\nTime difference (ms): {round(time_diff * 1000, 3)}')
206     print(f'\nTime difference (%): {round(time_diff_percent, 3)}%')
```

terminal output. It displays the elapsed time for both index construction techniques, as well as the time difference in ms and %, as follows.

```
Retrieving tokens ...
Tokens retrieved!
Collection size: 21578
Building SPIMI index ...
SPIMI index built!
Building Naive index ...
Naive index built!

===== STATISTICS =====

Token streams creation time (sec): 10.812

-----
| SPIMI Construction |
-----

Time taken for SPIMI index to process 10 000 terms (ms): 4.77
Size of SPIMI index: 2769 terms

-----
| Naive Construction |
-----

Time taken for Naive index to process 10 000 terms (ms): 6.836
Size of Naive index: 2769 terms

-----
| Differences |
-----

Time difference (ms): 2.066
Time difference (%): 30.224%
```


Subproject 2: First-Run Walkthrough

Running *subProject2.py* will trigger the *main()* function call and result in the following sequence of operations until terminal output.

Main

The main function below executes, retrieving the collection, building both indexes, and calling the *queryManager()* to handle continuous prompts for queries. The function begins by calling the same three initial functions as in the previous subproject. It calls *getTokens()*, as well as *spimi()* and *naïve()*, however it does not set thresholds for the indexes, and neither does it bother retrieving their construction times. Since these have all been seen in subproject1, we will skip to the last function call to *queryManager()*

```
397 def main():
398     """
399     The main function to initiate the search engine application. It retrieves token streams, builds indexes,
400     and launches the query manager to handle user queries.
401
402     This function is the starting point of the application.
403     """
404     tokenStreams = getTokens()
405     spimiIndex, _ = spimi(tokenStreams)
406     naiveIndex, _ = naive(tokenStreams)
407     queryManager(spimiIndex, naiveIndex, tokenStreams)
```

queryManager()

One of the main functions which handles all functionality to this subproject is the *queryManager()*, which is seen below. This function is responsible for continuously prompting the user for input parameters to the auxiliary functions necessary to the overall functioning of subproject2. It continuously prompts the user for a query, until 'q' is entered, which terminates the engine. For a single term query, the function would skip the prompt for operations, but in the case of a multi-term query the first function call would be to *getOperation()*.

```
371 def queryManager(spimiIndex, naiveIndex, collection):
372     """
373     Manages the querying process. It prompts the user to enter queries and handles the retrieval
374     and display of results using the chosen index and ranking method.
375
376     Args:
377         spimiIndex (dict): The SPIMI index.
378         naiveIndex (dict): The Naive index.
379         collection (list): The document collection, used for BM25 ranking.
380
381     This function continues to prompt for queries until the user chooses to quit.
382     """
383     print('\n===== WELCOME TO THE REUTERS21578 SEARCH ENGINE =====\n')
384     print('\n** Enter \'q\' to quit **\n')
385     print('\n** Queries do not require boolean operators **\n')
386     while True:
387         query = input('Enter your query: ')
388         if query == 'q':
389             break
390         numKeywords = len(tokenizer(query))
391         operation = getOperation() if numKeywords > 1 else None
392         spimi, naive = getIndexes(spimiIndex, naiveIndex)
393         queryTermRanking, bm25Ranking = getRanking(operation, spimi)
394         queryTest(query=query, spimiIndex=spimi, naiveIndex=naive, collection=collection, operation=operation, queryTermRanking=queryTermRanking, bm25Ranking=bm25Ranking)
395         print('\n===== SEE YOU LATER :P =====\n')
```

getOperation()

This function, as seen below, is simply responsible for continuously prompting the user for a Boolean operation (AND, OR) until a valid input is received.

```

294 def getOperation():
295     """
296     Prompts the user to choose a boolean operation ('AND' or 'OR') for the query processing.
297
298     Returns:
299     | str: The chosen boolean operation ('AND' or 'OR').
300     """
301     valid = False
302     while not valid:
303         operation = input('Enter your operation (\AND\, \OR\): ')
304         # Check if the entered operation is valid
305         if operation in ['AND', 'OR']:
306             valid = True
307         else:
308             print(f'Sorry, {operation} is not a valid operation.')
309     return operation

```

Once back in queryManager() with an operation retrieved, the next function call is to getIndexes() which can be seen below. This function works the same way in principle, prompting the user until a valid input is entered. The user could select to query either index, or both. The function will return the required, pre-built indexes depending on the selection.

```

311 def getIndexes(spimiIndex, naiveIndex):
312     """
313     Prompts the user to choose which index(es) to query - SPIMI, Naive, or both.
314
315     Args:
316     | spimiIndex (dict): The SPIMI index.
317     | naiveIndex (dict): The Naive index.
318
319     Returns:
320     | tuple: A tuple of the selected indexes. None is used for unselected indexes.
321     """
322     valid = False
323     while not valid:
324         answer = input('Which index would you like to query? \n' = naive, \s' = spimi, \b' = both : ')
325         # Validate the user input and return the appropriate indexes
326         if answer in ['n', 's', 'b']:
327             valid = True
328         else:
329             print(f'Sorry, {answer} is not a valid answer.')
330     # Return the appropriate combination of indexes based on user choice
331     if answer == 's':
332         return spimiIndex, None
333     elif answer == 'n':
334         return None, naiveIndex
335     elif answer == 'b':
336         return spimiIndex, naiveIndex

```

Once back in queryManager(), with the selected indexes, the next function call is to getRanking(), which principally works the same as the previous two functions. It prompts the user until it receives a valid ranking selection. The user can select query term ranking if they previously selected their operation as *OR*. They can select bm25 ranking if at least one of their selected indexes is the spimi index, or they can input *None* for no ranking. The function returns Boolean flags for each ranking technique.

```

338 def getRanking(operation, spimiIndex):
339     """
340     Prompts the user to choose a ranking method for the query results.
341
342     Args:
343         operation (str): The chosen boolean operation ('AND' or 'OR').
344         spimiIndex (dict): The SPIMI index.
345
346     Returns:
347         tuple: A tuple (queryTermRank, bm25) indicating whether query term ranking or BM25 should be used.
348     """
349     queryTermRank = False
350     bm25 = False
351     valid = False
352     while not valid:
353         answer = input('Which ranking would you like? \'q\' = query term ranking, \'b\' = bm25, \'n\' = none : ')
354         # Validate the user input and set the ranking method accordingly
355         if answer in ['q', 'b', 'n']:
356             if answer == 'q' and operation != 'OR':
357                 print('Sorry, query term ranking is only for multi-term OR queries.')
358             elif answer == 'b' and not spimiIndex:
359                 print('Sorry, bm25 ranking is only available for the spimi index.')
360             else:
361                 valid = True
362                 if answer == 'q':
363                     queryTermRank = True
364                 elif answer == 'b':
365                     bm25 = True
366             else:
367                 print(f'Sorry, {answer} is not a valid answer.')
368
369     return queryTermRank, bm25

```

queryTest()

Once back in `queryManager()`, now with the ranking flags set, `queryManager()` makes its last function call of the current iteration to `queryTest()`. The `queryTest()` function as shown below, is the function responsible for processing the queries depending on the inputted runtime parameters, which is why it is the longest. However it works rather simply. It begins by calling the `tokenizer()` on the query to retrieve a list of the tokenized query terms. If the list is empty then that means the inputted query was invalid. If it has a length of 1 then that means it is a single term query, and if more than 1 it is a multi-term query. The function begins by verifying if the naïve index has been initialized, if so, it steps into the if block and verifies the nature of the query. If it's a single term query, it queries the naïve index for the query term and outputs the

```

8 def queryTest(query, spimiIndex=None, naiveIndex=None, collection=None, operation="OR", queryTermRanking=False, bm25Ranking=False):
9     """
10     Processes a search query against provided SPIMI and/or Naive indexes. It supports single or multi-term
11     queries with AND/OR operations and optional ranking (Query Term Ranking or BM25).
12
13     Args:
14         query (str): The search query.
15         spimiIndex (dict): The SPIMI index (optional).
16         naiveIndex (dict): The Naive index (optional).
17         collection (list): The document collection (used for BM25 ranking).
18         operation (str): The boolean operation to apply ('AND', 'OR').
19         queryTermRanking (bool): Flag to use Query Term Ranking.
20         bm25Ranking (bool): Flag to use BM25 ranking.
21
22     The function tokenizes the query, checks the appropriate index for each term, and then applies the
23     specified operation (AND/OR) to compute the final result set. It also handles ranking if specified.
24     """
25     # Tokenize the input query
26     queryTerms = tokenizer(query)
27
28     # Check if the query is empty after tokenization
29     if len(queryTerms) == 0:
30         print(f'Sorry, \"{query}\" is not a valid query')
31         return
32
33     # Process with Naive Index if available
34     if naiveIndex:
35         print("\n===== NAIVE INDEX =====\n")
36         # Process single-term queries
37         if len(queryTerms) == 1:
38             result = naiveIndex.get(queryTerms[0]) # Get postings for the single query term.
39         else:
40             # Process multi-term queries
41             postingsLists = []
42             result = None
43             for term in queryTerms:
44                 postingsList = naiveIndex.get(term) # Get postings for each query term.
45                 postingsLists.append(postingsList)
46             # Apply boolean operations
47             if operation == 'AND':
48                 result = conjunction(postingsLists)
49             elif operation == 'OR':
50                 result = disjunction(postingsLists, queryTermRanking)
51
52     # Print the results from Naive Index
53     if result:
54         print(f"QUERY: \"{query}\" OPERATION: {operation} RANKING: {'Query Term Ranking' if queryTermRanking else 'None'} POSTINGS LIST SIZE: {len(result)} POSTINGS LIST: {str(result)}\n")
55     else:
56         print(f"QUERY: \"{query}\" - Your query returned no results with operation: {operation}.\n")
57

```

result. In the case of a multi-term query, the function iterates through the query terms, and retrieves their postings lists.

conjunction()

It then calls either conjunction() or disjunction() depending on the operation chosen at runtime. If AND, then conjunction is called.

```
189 def conjunction(postingsLists):
190     """
191     Performs an AND operation on a list of postings lists, returning the intersection of these lists.
192
193     Args:
194     | postingsLists (list of lists): A list where each element is a postings list (list of docIDs).
195
196     Returns:
197     | list: The intersection of the postings lists, representing documents that contain all terms.
198     """
199     # Return None if any of the postings lists is None
200     if None in postingsLists:
201         return None
202
203     # Sort postings lists by length for efficient intersection
204     sortedPostingsLists = sorted(postingsLists, key=len)
205     # Start with the shortest postings list
206     finalPostingsList = sortedPostingsLists[0]
207
208     # Intersect with each subsequent postings list
209     for postingsList in sortedPostingsLists[1:]:
210         finalPostingsList = intersect(finalPostingsList, postingsList)
211         # If intersection is empty, no further processing is needed
212         if not finalPostingsList:
213             break
214     return finalPostingsList
```

Intersect()

The conjunction function calls the intersect() function on each of the postings lists that have been passed as a parameter beginning with the two smallest. It does so two at a time, and combines the results in order. The intersection function can be seen below, which functions as a regular intersection, combining postings lists by common docIDs while avoiding duplicates and incrementing with respect to the list positioned at the smallest docID, until the end of one of the lists has been reached.

```
248 def intersect(pList1, pList2):
249     """
250     Computes the intersection of two postings lists.
251
252     Args:
253     | pList1 (list): The first postings list.
254     | pList2 (list): The second postings list.
255
256     Returns:
257     | list: The intersection of pList1 and pList2.
258     """
259     result = []
260     i = j = 0
261     # Iterate through both lists to find common elements
262     while i < len(pList1) and j < len(pList2):
263         if pList1[i] == pList2[j]:
264             result.append(pList1[i])
265             i += 1
266             j += 1
267         elif pList1[i] < pList2[j]:
268             i += 1
269         else:
270             j += 1
271     return result
```

Back in the queryTest() function, in the case of a conjunction, the resulting list of docIDs is returned and displayed.

disjunction()

However, if the selected operation is *OR* then the `disjunction()` function is called which can be seen below. It simply unions the postings lists in order and without duplicates, unless they are all *None*.

```
216 def disjunction(postingsLists, queryTermRanking=False):
217     """
218     Performs an OR operation on a list of postings lists, returning the union of these lists. Can also
219     apply query term ranking based on term frequency across the postings lists.
220
221     Args:
222         postingsLists (list of lists): A list where each element is a postings list (list of docIDs).
223         queryTermRanking (bool): Flag to apply query term ranking.
224
225     Returns:
226         list: The union of the postings lists, with optional query term ranking applied.
227     """
228     # Filter out None postings lists (terms not found)
229     validPostingsLists = [plist for plist in postingsLists if plist is not None]
230
231     # If all postings lists are None, return None
232     if not validPostingsLists:
233         return None
234
235     # Flatten the list of lists into a single list using itertools.chain
236     unionPostings = list(chain(*validPostingsLists))
237
238     # If query term ranking is not applied, remove duplicates and sort
239     if not queryTermRanking:
240         finalPostings = sorted(set(unionPostings))
241     else:
242         # Sort by document ID and apply query term ranking
243         finalPostings = sorted(unionPostings)
244         finalPostings = queryTermRank(finalPostings)
245
246     return finalPostings
```

queryTermRank()

In the case of `queryTermRanking` flag is set to true, then this is handled in the `disjunction()` function by calling the `queryTermRank()` function on the resulting disjunct set. The `queryTermRank()` can be seen below, where it maintains a local dictionary of docIDs, incrementing each of their counts everytime they are encountered in the `postingsList`. It then sorts the dictionary with respect to value, and returns the dictionary as a list of ordered docID-count tuples.

```
273 def queryTermRank(postingsList):
274     """
275     Applies query term ranking to a postings list. Ranks documents based on the number of query terms they contain.
276
277     Args:
278         postingsList (list): A postings list (list of docIDs).
279
280     Returns:
281         list: A list of tuples (docID, score) where score is the count of query terms in the document.
282     """
283     # Count the frequency of each document ID in the postings list
284     docFreq = {}
285     for docID in postingsList:
286         docFreq[docID] = docFreq.get(docID, 0) + 1
287
288     # Sort documents by frequency (score) in descending order
289     rankedDocs = sorted(docFreq.items(), key=lambda item: item[1], reverse=True)
290     return [(docID, score) for docID, score in rankedDocs]
```

Back in `queryTest()`, now with the resulting disjunct set of docIDs in the case of an *OR* operation, the results are displayed to terminal. The function then breaks from the *if naïve* if block, and checks if the `spimi` index parameter was initialized. If yes, then it enters the if block and goes through the same steps as above for the naïve if block. **The only difference**, is two main points.

convertPostingsLists()

Firstly, since the spimi index's postings lists are lists of tuples and not regular docIDs, it calls a function to normalize the postings lists to a simple list of docIDs. The function can be seen below. It does so before calling either conjunction() or disjunction().

```
175 def convertPostingsLists(postingsLists):
176     """
177     Convert a list of postings lists with (docID, tf) tuples to a list of lists of docIDs.
178     If a postings list is None, it remains None.
179
180     Args:
181     | postingsLists (list): A list of postings lists, where each postings list contains tuples of (docID, tf).
182
183     Returns:
184     | list: A list of lists, where each inner list contains docIDs or is None.
185     """
186     # Convert each postings list to a list of docIDs, or keep as None if postings list is None
187     return [[docID for docID, tf in postingsList] if postingsList is not None else None for postingsList in postingsLists]
```

bm25()

Secondly, the only other difference is that it checks if the bm25 flag has been set. If yes, it makes a function call to the bm25() function below. It maintains the scores of documents in a local dictionary rankedResults. It does so by retrieving and iterating over the docIDs in each of the queryTerms' postings lists. When the docID matches one of the docIDs in the result parameter, then it computes the score, and stores/updates it in the rankedResults dictionary. Once all queryTerms have been iterated over the dictionary is ordered by score and returned as a sorted list of docID-score tuples.

```
126 def bm25(queryTerms, result, index, collection, k1, b):
127     """
128     Applies the BM25 ranking formula to the result set of a query. It calculates scores for documents
129     based on term frequency, document frequency, and document length.
130
131     Args:
132     | queryTerms (list): List of tokenized query terms.
133     | result (list): List of document IDs obtained from the initial search.
134     | index (dict): The search index.
135     | collection (list): The document collection, to compute document lengths.
136     | k1 (float): BM25 parameter k1.
137     | b (float): BM25 parameter b.
138
139     Returns:
140     | list: A sorted list of tuples (docID, score), ranked according to BM25.
141     """
142     N = len(collection) # Total number of documents
143     L_total = sum(len(doc) for doc in collection) # Total length of all documents
144     L_avg = L_total / N # Average document length
145
146     rankedResults = {}
147
148     if result is None:
149         return None
150
151     # Calculate BM25 score for each term in each document
152     for term in queryTerms:
153         postingsList = index.get(term, [])
154         df = len(postingsList) # Document frequency
155         idf = math.log(N / df) if df != 0 else 0 # Inverse document frequency
156
157         for docID, tf in postingsList:
158             if docID and docID in result:
159                 L_d = len(collection[int(docID)]) # Document length for the current docID
160                 # BM25 formula
161                 score = idf * ((tf + (k1 + 1)) / (k1 * ((1 - b) + b * (L_d / L_avg)) + tf))
162                 # Accumulate score for each document
163                 if docID not in rankedResults:
164                     rankedResults[docID] = 0.0
165                 rankedResults[docID] += score
166
167     # Check if all scores are zero
168     if all(score == 0.0 for score in rankedResults.values()):
169         return None
170
171     # Sort the results by score in descending order
172     sortedRankedList = sorted(rankedResults.items(), key=lambda item: item[1], reverse=True)
173     return sortedRankedList
```

Once this is complete, the queryManager() then prompts the user for another query, restarting the whole process.

***** For output logs, please refer to Appendix A and Appendix B of 4013360_report.pdf**